

```

1 #=====
2 # Welcome to R! This is a primer to demystify R a bit and help get you on your feet.
3 #=====
4 # Any questions, contact Ty Taylor: tytaylor@email.arizona.edu
5
6 # << INTRODUCTION >> -----
7
8 # This primer aims to teach you the basics of using R Studio, and speaking the R language.
9 # The tutorial is EASY, you are just reading instructions and typing code. In each short
10 # section, different aspects of R are explained, and shown by example.
11
12 # << R overview >> -----
13
14 # R reads your data (e.g., from an excel or text file) and makes an image of it, which is
15 # stored in temporary memory. When you edit your data, you are only editing an image of
16 # it; the original dataset is unaffected. You can then create graphs and altered
17 # datasets, which can be written back to a permanent file. This is called 'nondestructive
18 # editing'. You can also create objects (like matrices) from directly within R, as opposed
19 # to reading it in from a file.
20
21 # R can do almost anything you can think of doing with your data, which greatly expands
22 # the amount of creativity you can apply to your analyses. By using scripts, you keep a
23 # record of everything you've done to manipulate and analyze your data. Those benefit you
24 # later, and can be published or shared with colleagues.
25
26 # << Primer tutorial instructions >> -----
27
28 # There are two scripts (.R files), one with comments+code, and one with comments only
29 # (with BLANK in the title).
30
31 # There is a data file included with these scripts called "glopnet.csv". This is a
32 # publicly available dataset containing trait data for numerous plant species.
33
34 # First, make sure the data file (glopnet.csv) and the two scripts are in the same folder.
35
36 # >> Next, download R and R Studio (two separate programs, download R first, then R
37 # Studio).
38
39 # Open R Studio. Open both of the .R files in R Studio (File -> Open -> navigate to script
40 # files).
41
42 # You'll be working in R Studio with the file with "BLANK" in the title. The other will be
43 # open just for reference (and to demonstrate how R Studio can have multiple scripts open
44 # at a time, each in a different tab).
45
46 # << In class >> -----
47
48 # Just follow along in the BLANK script, and type the code that I type beneath the
49 # comments.
50
51 # << On your own >> -----
52
53 # Open the PDF version of the comments+code script and view it along side R Studio. It
54 # helps if you can view this on a separate screen, or print it.
55
56 # Wherever a line does NOT have a hash # symbol, that is a line of code for you to type
57 # into the appropriate place in the BLANK script in R Studio. In the PDF version, code is
58 # highlighted. I recommend actually typing in the code and not copying and pasting. The
59 # point here is to develop a working knowledge of scripting by actually doing it.
60
61 # Execute each line of code IN ORDER. (instructions for executing code are in Section 1
62 # and 2).
63
64 # Note that capital letters make a difference in your code, so be sure to copy code

```

```

65 # exactly. Also, in this tutorial you will be creating objects in R and using them later,
66 # so if you jump to a later section, the code might not work if it calls an object that
67 # should have been created earlier.
68
69 # **If things aren't working, you can always highlight and run all the code in the
70 # completed script up to your current line. That way all objects created from previous
71 # lines will be recreated correctly.
72
73
74 #=====
75 # Tutorial contents
76 #=====
77
78 # Part 1: How easy are ANOVAs, regressions, and plots?
79 # Part 2: Getting to know R Studio and basic R functionality
80 # Part 3: Objects, object classes, text
81 # Part 4: Vectors and vector indexing
82 # Part 5: Built-in functions and R Help
83 # Part 6: Matrices - creating, indexing, and manipulating
84 # Part 7: Data frames - R's most common data format
85 # Part 8: Importing data as data frames
86 # Part 9: Getting info from bigger data frames
87 # Part 10: Packages
88 # Part 11: Basic troubleshooting
89
90
91 #=====
92 # Part 1: How easy are ANOVAs, regressions, and plots?
93 #=====
94
95 # ** NOTE ** This section is a jump ahead. It is just to show you how easy it can be to
96 # use R to do ANOVAs, regressions, and plots. Many aspects of what you see in these first
97 # few lines will be explained in the later sections. Just follow the instructions below to
98 # execute the lines one by one and see what happens! Or, if you want the practice, delete
99 # the lines from the BLANK version and type them in following the completed version.
100
101 # >> FIRST: Setting your working directory and executing code from the script
102
103 # 1) First, you need to set the 'working directory' (first line of code below), DO NOT
104 # type the line of code you see in the completed script. Instead, in the drop-down menu at
105 # the top of R Studio, go to Session -> Set working directory -> To source file location.
106 # If you want to be able to smoothly execute this whole script later without doing that
107 # manual step, copy the line of code it has printed in the console below (starting with
108 # "setwd()") and paste it into the script below.
109
110 # 2) After you type a line of code into your script in R Studio (copying the completed
111 # script) just hit Ctrl+ENTER (Command+ENTER for mac) with the blinking text cursor active
112 # on that line. That is how lines of code are executed from the script window (all this
113 # explained in the following sections).
114
115 #--Set working directory to source file location (see instructions above).
116 setwd("~/Directory path to the folder containing the R Primer script")
117
118 #--Read in glopnet.csv as a data frame.
119 glop <- read.table (file="glopnet.csv", sep=",", header=TRUE)
120
121 #--Have R show you a list of the column names.
122 names(glop)
123
124 #--Do an ANOVA to see if log(leaf life span) varies among genera.
125 # "Pr(>F)" is your p value.
126 log.LL.aov <- aov (log.LL ~ Genus, data=glop)
127 summary (log.LL.aov)
128

```

```

129 #--Make a scatter plot of log(leaf mass per area) with log(leaf life span)
130 plot (glop$log.LMA, glop$log.LL)
131
132 #--Do a linear regression analysis to see if the correlation between LMA and LL is
133 # significant. Pr(>|t|) for glop$log.LL is your p value for the correlation.
134 LL.LMA.reg <- lm (glop$log.LL ~ glop$log.LMA)
135 summary (LL.LMA.reg)
136
137 #--Add the regression line in blue to your scatter plot.
138 abline (LL.LMA.reg, col="blue")
139
140
141 #=====
142 # Part 2: Getting to know R Studio and basic R functionality
143 #=====
144
145 #-----
146 # R Studio layout
147 #-----
148
149 # R studio is just a nice visual interface with R. You can open up R on its own to see
150 # what the default visual interface looks like.
151
152 # R Studio has 4 window panes: Script window (upper left), console (below), two
153 # multi-purpose windows (right). You can drag the edges to adjust sizes.
154
155 # NOTE: What is a script?:
156 # A script is like a story about your data analysis. It is written with comments in plain
157 # English, and code in R language. R reads the story, ignoring the English, and makes it a
158 # reality. You can execute code directly in the console. But writing scripts keeps a
159 # record of every step of your analysis, and allows you re-run the analysis later (maybe
160 # on an altered dataset), or change steps in the analysis.
161
162 # Script window: This is where you develop scripts. Some nice automatic functions include
163 # color coding, indenting, and bracket/quote closing. Multiple scripts can be stored in
164 # separate tabs that can be selected above. It's useful to keep several old scripts open
165 # in these tabs to use as a reference while you work.
166
167 # Console window: This is like the regular R console window. This is where your scripts
168 # run, and where visualized results are returned (but if you don't ask R to show you
169 # something, it won't).
170
171 # Bottom right window: This is mostly used for viewing Help files and Plots you've created
172 # during a session.
173
174 # Upper right window: Here you can view a list of items in your 'Environment'. These are
175 # all the objects you've created or loaded into your environment during this session.
176
177 #-----
178 # Basic functionality
179 #-----
180
181 # >> Executing code from the script window; examples with basic math functions:
182 # Type in each line of code below the comment. Execute each line by having your active
183 # cursor on that line, and hitting Ctrl+ENTER (or Command+ENTER). Notice the executed line
184 # of script is shown in blue in the console window, and the result (if your line asks for
185 # one) is shown in black. You do not need to type in the "#" symbols or the words
186 # following them. Those are just comments (explained below).
187
188 #--addition
189 2+2
190 #--multiplication
191 2*2
192 #--division

```

```

193 2/2
194 #--exponent
195 2^3
196
197 # >> Executing only part of a line, or multiple lines from the script window:
198 # Highlight the desired section of script and hit Ctrl+ENTER (Command+ENTER).
199
200 #--Try executing just the "2*2" part of this line. The code and result are shown in the
201 # console.
202
203 4+9+2*2
204
205 #--Now try executing two of the above lines of code at once.
206
207 # >> Executing lines from the console:
208 # Just click inside the console window, type your code, and hit ENTER.
209 # NOTE: if you accidentally hit command (or ctrl) ENTER, it will execute instead the line of
210 # code selected in your script window. This can be really annoying.
211
212 # >> Toggling through previous lines of code in the console:
213 # Click inside the console window. Use the up and down arrows to toggle through
214 # previously executed lines of code. You can then edit them, then hit ENTER to
215 # execute.
216
217 # >> Commenting:
218 # Using the "#" symbol tells R to disregard anything following that symbol. Use it to
219 # nullify a line of code or to make comments. Execute the 2*2 line below, and notice that
220 # in the console, the answer is not shown because the formula was disregarded.
221
222 # 2*2
223
224 # >> Spaces:
225 # R disregards spaces between items. USE THEM to make your code easier to read.
226
227 (4+4) / (2+2)
228
229
230 #=====
231 # Part 3: Objects, the equivalence principle, object classes, text.
232 #=====
233
234 # Programming relies heavily on the creation and manipulation of objects. Objects have
235 # names, and contain structured information (data). You can either manipulate the object
236 # itself, its sub-objects, or the data the object contains. This will become clear.
237
238 #-----
239 # Creating, viewing, and using objects; the equivalency principle
240 #-----
241
242 # >> Defining an object:
243 # An object is defined using the "=" symbol, or an arrow "<=". The object is then stored
244 # in memory (but not saved in any file).
245
246 #--Make the object 'a' equal to 20.
247 a <- 20
248
249 # NOTE: The "=" and "<=" are interchangeable, but I like to use the arrow when defining an
250 # object, and the = when assigning values to cells. It's easier to scan your script for
251 # objects that way.
252
253 # Notice that the line of code (in blue) is shown in the console, but no answer in black.
254 # That's because you didn't ask R to show you anything, you just told it that 'a' equals
255 # 20.
256

```

```

257 # Notice also that 'a' now appears in your Workspace window at the upper right of your
258 # screen.
259
260 # >> Viewing an object:
261 # Just type the object's name and execute.
262 a
263
264 # >> Using objects:
265 # Just add them in your code. If a = 20, then a+1 = 21.
266 a + 20
267
268 # >> The equivalency principle **VERY IMPORTANT**:
269 # The above example shows how the object 'a' and the number 20 are now essentially
270 # equivalent, although an object can have extra features like names. Anything that you
271 # could do with 20, you can do with 'a', so consider an object and its content essentially
272 # interchangeable. You will have to remind yourself of this in less obvious situations.
273
274 #--Make an object 'b' equal 30, and ask R what a plus b is.
275 b <- 30
276 a + b
277
278 # >> Naming objects:
279 # Objects can have more complicated names with periods and underscores, just don't include
280 # mathematical operators like "-", and NO SPACES IN NAMES.
281 my_new.object <- 10
282
283 #-----
284 # Object classes
285 #-----
286
287 # There are different types, or 'classes', of objects in R. The class of the object
288 # determines what kinds of things can be done with it, what types of data structures are
289 # allowed, etc.
290
291 #--Use the class() function to determine what class of object 'a' is. (The syntax of
292 # functions will be explained later.)
293 class(a)
294
295 # 'a' is a numeric object.
296
297 #-----
298 # Text strings (class Character)
299 #-----
300
301 # >> Defining a text string:
302 # Text is defined by 'single' or "double" quotes. A "string" is a series of characters.
303 text.1 <- "this is a character string"
304 text.1
305
306 #--Ask R what the class of 'text.1' is.
307 class(text.1)
308
309 # Because the content of your object was enclosed in "quotes", R automatically assigned it
310 # the class 'character'.
311
312
313 #=====
314 # Part 4: Vectors and vector indexing
315 #=====
316
317 #-----
318 # Vectors
319 #-----
320

```

```

321 # Vectors are linear groups of things of the same class (e.g., numeric, character). These
322 # enable you to assign a whole group of things to an object. A vector can have a length of
323 # 1, like our object 'a'. Longer vectors are created by the function c() for "concatenate".
324 # Items are separated by commas. Math can go in between commas.
325 vector.one <- c(1,2,2+2)
326 vector.one
327
328 # A number series is denoted by ":".
329 vector.two <- 1:9
330 vector.two
331
332 #--You can do math with vectors.
333 vector.two * 2
334 vector.two * vector.one
335
336 # ** NOTE: When you multiply one vector by another, the values in corresponding positions
337 # are multiplied. This is not the same as matrix-multiplication of vectors. You do that by
338 # %*% instead of just *.
339
340 #--You can combine vectors by making a vector of vectors.
341 vector.three <- c(vector.one, vector.two)
342 vector.three
343
344 #--Remember the equivalency principle. When building a vector, any code that results in a
345 # single item, or a vector of items, can be inserted between commas when defining the
346 # vector.
347 vector.four <- c(1, 2, vector.two*2, 7:10, c(1,2,3))
348 vector.four
349
350 #--Remember that vectors must contain elements of the same class! When you combine
351 # multiple classes, unexpected things happen.
352 vector.five <- c(1,"a",2)
353 vector.five
354 class(vector.five)
355
356 # R reduces mixed classes in a vector to a single class according to its hierarchy of
357 # classes.
358
359 #-----
360 # Indexing vectors
361 #-----
362
363 # 'Indexing' is pointing to specific items or groups of items in an object. INDEXING IS
364 # THE SINGLE MOST IMPORTANT SKILL IN R. Don't forget that!
365
366 # One-dimensional indices, for one-dimensional vectors, are denoted by [ ] following the
367 # object name.
368
369 #--Ask R what the third element in vector.one is.
370 vector.one[3]
371
372 #--Make an object 'prod' equal to the product of element 3 from vector.one and element 2
373 # from vector.two. Ask R to show you 'prod'.
374 prod <- vector.one[3] * vector.two[2]
375 prod
376
377 #--Give vectors one and two names that are faster to type.
378 v1 <- vector.one
379 v2 <- vector.two
380
381 # NOTE: this does not actually change their names, it just makes new objects that are
382 # equivalent but have shorter names.
383
384 # You can specify multiple indices by inserting a vector in [ ]

```

```

385
386 #--Return the 1st and 3rd elements of v1.
387 v1 [ c(1,3) ]
388
389 #=====
390 # Part 5: Built-in functions, R Help, and external help
391 #=====
392
393 #-----
394 # Built-in functions and R Help
395 #-----
396
397 # << R is functions >> -----
398
399 # There are many built in, and you can make your own. Functions can do things to numbers,
400 # text, whole datasets, everything really.
401
402 #--Functions for basic stats: mean(x), sd(x), max(x), min(x)
403 mean (v2)
404 sd (v2)
405 max (v2)
406 min (v2)
407
408 # << Function structure >> -----
409
410 # All functions look like: function_name (argument 1, argument 2, ...)
411 # The name precedes parentheses, which enclose arguments, which are separated by commas.
412
413 #--Use the function rnorm() to generate 10 numbers drawn from a normal distribution with
414 # mean 5 and standard deviation 2.
415 rnorm (n=10, mean=5, sd=2)
416
417 # << Arguments >> -----
418
419 # Think of 'arguments' as questions needing answers. The arguments of a function have
420 # names, and occur in a particular order. In rnorm(), 'n', 'mean', and 'sd' are arguments,
421 # which occur in that order.
422
423 # >> Calling the arguments without naming them:
424 # If you call the arguments in order, you don't need to name them.
425 rnorm (10, 5, 2)
426
427 # >> Calling the arguments out of order:
428 # If you name the arguments, you don't need to call them in order.
429 rnorm (mean=5, n=10, sd=2)
430
431 # >> Argument defaults:
432 # Some arguments have default values, and therefore don't need to be specified by you if you
433 # don't want to change them. rnorm() has defaults of mean=0 and sd=1.
434 rnorm (n=10)
435
436 # If you fail to specify an input for an argument with no default, R will return an error.
437 rnorm ()
438
439 # << R Help >> -----
440
441 # In order to use a function, you must know what the arguments (questions) are, and what
442 # kind of inputs (answers) are allowed. For this, start with R Help.
443
444 # To get documentation on a function, type a question mark followed by the function name,
445 # and execute: ?function_name
446
447 # Functions are defined in the R Help tab in the lower right window. Help will tell you
448 # which arguments can or must be defined, and what their defaults are.

```

```

449
450 #--Get help on the paste() function. paste() concatenates character strings.
451 ?paste
452
453 # >> Structure of the R Help file:
454 # Under "Usage", the order, and default settings for the arguments are shown.
455 # - Note the default input for the separator argument of paste, 'sep', is equal to a space
456 # " ".
457 # Under "Arguments", the arguments are defined. Read them.
458 # The detail of information increases as you scroll down, and there are examples at the
459 # bottom (but they are often more complicated than they need to be).
460
461 #-----
462 # Practice with the paste() and rep() functions.
463 #-----
464
465 # >> The paste() function:
466 # The first arguments of paste() are the text strings to be concatenated, each separated
467 # by a comma. The separator argument 'sep' can be specified, or left to its default " ".
468
469 #--Make a new text string and paste it to text.1 (which we created earlier).
470 text.2 <- "with some more text"
471 paste(text.1, text.2)
472
473 #--Now do that again, but make the separator be "_" instead of the default " ".
474 paste(text.1, text.2, sep="_")
475
476 #--Set the separator to none, or no space.
477 paste(text.1, text.2, sep="")
478
479 # >> rep() repeats things
480 ?rep
481
482 #--Repeat the letter a 20 times.
483 rep("a", 20)
484
485 # POP QUIZ 1: What happens if you DON'T put the quotes around "a"? Why did that happen?
486 # - Pop quiz answers at bottom of tutorial.
487
488 # >> Nesting functions:
489 # There is no problem with putting functions inside of other functions.
490
491 #--Make a vector 'v4' with five 10's in a row, followed by a 4.
492 v4 <- c(rep(10, 5), 4)
493 v4
494
495 # Note how we embedded the function rep() in the function c() defining the vector. Also
496 # note how you need to start being very careful with your placement of parentheses and
497 # commas!
498
499 # Notice the equivalency principle in use above:
500 # If it's ok to write: c(10, 10, 10, 10, 10, 4)
501 # And rep(10,5) is equal to 10, 10, 10, 10, 10
502 # Then we can write: c(rep(10,5), 4)
503
504 #-----
505 # External help
506 #-----
507
508 # How do you know which function to use if you've never seen it before?
509 # Use Google! Seriously. Try searching "concatenate text in r".
510
511 # **Stack Overflow** site: Of all of the sites providing answers to your Google search
512 # for coding questions, Stackoverflow will almost always have the best answers in the most

```



```

513 # readable format. I ALWAYS look there first.
514
515 # You can also search help topics within the R help window.
516 # Or type, e.g., ??concatenate in the console to search the help pages.
517
518
519 #=====
520 # Part 6: Matrices - creating, indexing, and manipulating
521 #=====
522
523 #-----
524 # Creating, indexing, and manipulating matrices
525 #-----
526
527 # Matrices are the two-dimensional version of a vector. Like vectors, they are limited to
528 # containing items all of the same class. (Can't mix numbers and character strings, for
529 # example.)
530
531 # >> To create a matrix, use the matrix() function:
532 # In the first argument (code before the first comma), put in a vector.
533 # Then add an argument defining the dimensions in terms of the number of rows or columns.
534 ?matrix
535
536 #--Make a 2-row matrix named 'mat1' out of the vector of numbers 1 through 6.
537 mat1 <- matrix (1:6, nrow=2)
538
539 #--View mat1
540 mat1
541
542 # NOTE: R will always build your matrix by columns first, then by rows! Remember that.
543
544 #--Math can be done on matrices.
545 # Multiply everything in the matrix by 2.
546 mat1 * 2
547
548 # NOTE: This is not true matrix multiplication! Regular mathematical operators are only
549 # applied element by element in order. For true matrix multiplication, use the %*%
550 # operator (e.g., mat1 %*% mat2).
551
552 # << Indexing matrices >> -----
553
554 # Remember, indexing means pointing at a desired element or group of elements.
555
556 # Indices for two dimensional objects are denoted by [ , ] after the object name. The ROW
557 # NUMBER goes LEFT of the comma, the COLUMN NUMBER goes RIGHT. Imprint that in your brain:
558 # [Row, Column]. [Row, Column]. [Row, Column].
559
560 #--What is the entry in the 2nd row and 3rd column of mat1?
561 mat1 [2,3]
562
563 # >> Call an entire column or row by leaving the other dimension index blank.
564
565 #--Return all the entries in the 2nd column of mat1, multiplied by 2.
566 mat1 [ ,2] * 2
567
568 #--Redefine the entries in the second column of mat1 such that they are multiplied by 2.
569 mat1 [ ,2] = mat1 [ ,2] * 2
570 mat1
571
572 #--Return columns 2 and 3 (using a vector inside of [ , ] to specify multiple columns).
573 mat1 [ ,2:3]
574
575 # >> Indexing matrices one-dimensionally:
576 # Matrices actually ARE vectors, with a two-dimension attribute. Therefore, you can call

```

```

577 # the ith element just as in a vector (i.e. without a comma). Remember, Matrices are
578 # populated and counted by columns first (top to bottom; left to right).
579
580 #--Ask R to show you elements 3 through 6 of mat1.
581 mat1 [3:6]
582
583
584 #=====
585 # Part 7: Data frames - R's most commonly used data format
586 #=====
587
588 # Data frames are like spreadsheets, so they will be the most familiar to you. They can
589 # have column and row names, and many different types of data entries in a matrix-like
590 # format.
591
592 # Data frames are composed of columns, all of the same length. Each column is a vector.
593 # Therefore all elements in a given column must be of the same class.
594
595 #--Create a data frame out of mat1, using the data.frame() function.
596 ?data.frame
597 df1 <- data.frame (mat1)
598 df1
599
600 # NOTE: now the [,] indices are replaced by variable (column) names and row names
601 # (numbers) when R shows you the data frame in the console.
602
603 #--Change the column names using the colnames() function.
604 ?colnames
605 colnames (df1) <- c("one","two","three")
606 df1
607
608 # >> The column names are a vector of names (of class 'character'), and can be indexed
609 # like a vector.
610
611 #--View the column names of your data frame. Note that they show up in quotes because they
612 # are character strings.
613 colnames (df1)
614
615 #--Change just the third column name by indexing the third element in the colnames of df1.
616 colnames (df1) [3] <- "col.3"
617 df1
618
619 # NOTE: I didn't have to write c("col.3") because it's just one item, not a group.
620
621 #-----
622 # Indexing and adding columns to data frames.
623 #-----
624
625 # >> Two-dimensional indexing [ , ]
626 # Data frames can be indexed like matrices with [ , ] following the data frame name, and
627 # entering row numbers left of the comma and column numbers right.
628
629 #--Call the second row of columns 2 and 3 of df1.
630 df1 [2, 2:3]
631
632 #--Notice the difference in what R returns from df1 compared to the same indices in mat1.
633 mat1 [2, 2:3]
634
635 #--Check the class() of the two different datasets returned.
636 class (df1 [2, 2:3])
637 class (mat1 [2, 2:3])
638
639 # A vector of numbers is returned from the matrix, and a data frame is returned from the
640 # data frame.

```

```

641
642 # >> $ indexing
643 # You can call a column by name using $, as in "data.frame$column.name".
644
645 #--Call the column named "two" of df1.
646 df1$two
647
648 # >> One-dimensional indexing in a column with [ ]
649 # Since a column is just a vector, you can index within a column as you would with a
650 # vector.
651
652 #--Call the 1st item in column "two" of df1.
653 df1$two [1]
654
655 # >> Indexing columns by name
656 # You can call a column by name inside of the [ , ] using quotes.
657
658 #--Call column "two" of df1.
659 df1 [ , "two"]
660
661 #--Call the second row item in column "two" using [ , ] with a named column index.
662 df1 [2, "two"]
663
664 # >> Adding a column
665 # Add a column to the data frame just by defining a new one with $.
666
667 #--Make a new column for df1 called "plot" with blank text string entries.
668 df1$plot <- ""
669 df1
670
671 # NOTE: Since you didn't specify how many blank entries, it just populated the whole
672 # column with them.
673
674 #--Make the two entries in df1$plot equal to "treatment" and "control".
675 df1$plot = c("treatment", "control")
676 df1
677
678 # NOTE: it was not necessary to first create a column with blank entries, the entries
679 # could be defined right away. Sometimes it makes things more clear to do them in multiple
680 # stages though.
681
682 #-----
683 # Getting basic information about your data frames and other objects
684 #-----
685
686 # >> str() returns the structure of your object.
687 # What is the structure of df1?
688 str (df1)
689
690 # NOTE: "plot" is of the class 'character' (chr), while the other columns are numeric.
691
692 # >> summary() gives you a statistical summary of the object.
693 summary (df1)
694
695 # >> nrow() asks how many rows the object has.
696 nrow (df1)
697
698 # >> ncol() asks how many columns the object has.
699 ncol (df1)
700
701 # >> length() gives you the length (number of elements) of a vector or list.
702 # Ask how many elements there are in df1$plot.
703 length (df1$plot)
704

```

```

705 #--How many elements are in df1?
706 length(df1)
707
708 # POP QUIZ 2: What unit is considered a single 'element' of df1? (Go back to the definition
709 # of data frames at the beginning of this Part.)
710 # - Pop quiz answers at bottom of tutorial.
711
712
713 #=====
714 # Part 8: Importing data as data frames
715 #=====
716
717 #-----
718 # The working directory
719 #-----
720
721 # Setting a 'working directory' simplifies importing and exporting. When you ask R to
722 # import (read) a file, or export (write) a file, it will point automatically to the
723 # working directory, so you don't have to type in a directory path.
724
725 # >> The easiest way: set wd to the script file location.
726 # In R Studio, go to "Session" in the drop-down menu at the very top.
727 # Click "Set working directory" -> "to source file location".
728 # This will point R to the folder containing THIS SCRIPT.
729 # Notice that the line of code for doing that is implemented in the console. To
730 # incorporate that line into your script (so you don't have to take that extra step in the
731 # future) just copy it from the console and paste it into the script.
732
733 setwd("~/Directory path to the folder containing the R Primer script")
734
735 # NOTE: you can use the command setwd() to set the working directory to any other file
736 # path you want. The "~" points to your Documents folder. The slashes must be forward /
737 # not back \.
738
739 #-----
740 # Reading in a file as a data frame.
741 #-----
742
743 # >> read.table() function
744 #--Read in the csv file "glopnet.csv":
745 # Make sure glopnet.csv is in the folder containing this script.
746 # Set your working directory "to source file location" as described above.
747 # Use the read.table() function to read in the file and assign it to an object called
748 # 'glop'.
749 ?read.table
750 glop <- read.table (file="glopnet.csv", sep=",", header=TRUE)
751 glop
752
753 # Note that when you view glop, it is a large data frame so will not show you all of it.
754 # We'll look at easy ways to see only parts of it and get specific information in the next
755 # section.
756
757 # >> read.csv() function
758 # This is a shortcut for '.csv' files. The default separator is "," and header=TRUE.
759 glop <- read.csv (file="glopnet.csv")
760
761 #-----
762 # Using custom directories, pasting the file name. (Preferred method!)
763 #-----
764
765 # Your file management will usually be neater if you work with multiple folders. E.g., you
766 # might want your scripts in one folder, your datasets in another, and your figures in
767 # another.
768

```

```

769 # In that case, it's handy to create objects that contain the directory path as a text
770 # string.
771
772 #--Make an object dat.dir that contains the directory path with your data as a text
773 # string.
774 # - To get your directory path, navigate to your data file in your file finder/explorer
775 # - right click to get properties/get info
776 # - copy the file directory path
777 # - Paste it between quotes to define dat.dir.
778 # - If the path contains backslashes "\" you need to change them to forward slashes "/".
779 # Just highlight that text in your script, hit Ctrl+F (Command+F). In the 'find' window
780 # that comes up at the top, select the box "In selection". Find \ and replace-all with /.
781 # - FINALLY, add one / to the end of the text string.
782
783 dat.dir <-" /Users/ttaylor/Documents/Stats and R course/Tutorials/Section 1/"
784
785 #--Read in glopnet.csv using read.csv().
786 # This time, make the file name equal to paste0(dat.dir, "glopnet.csv")
787 glop <- read.csv (file = paste0 (dat.dir, "glopnet.csv"))
788
789 # This bypasses the working directory because you've pointed to a directory in the file
790 # name.
791
792 # NOTE: paste0() is just like paste(), but with the default separator = "".
793
794
795 #=====
796 # Part 9: Getting info from bigger data frames
797 #=====
798
799 # R will not print all of a big data frame in the console. But there are a variety of ways
800 # to get the information you want.
801
802 # >> Viewing the data frame as a spreadsheet, like in Excel:
803 # Go to the upper right R Studio window, in the Workspace tab, click on glop. It will open
804 # a new script tab and show you the first 1000 rows of glop in spreadsheet format.
805
806 # >> str(), shown earlier, is great for big data frames.
807 str (glop)
808
809 # >> head() shows you the first few rows of data.
810 # **NOTE: I use this function more than any other! Any time I make a change, or pull in
811 # data, this shows me whether things are as I expect them to be.
812 head (glop)
813
814 # >> tail() shows you the last few rows of data.
815 tail (glop)
816
817 # >> names() gives you a list of column names in the data frame. (For a data frame, it's
818 # the same as colnames() )
819 names (glop)
820
821 # Remember, the column names are a vector, and can be indexed as such.
822
823 # >> unique() lists all the unique entries in an object.
824 # Ask R for all the unique entries in the "Dataset" column.
825 unique (glop$Dataset)
826
827 # >> length() tells you how many elements are in an object.
828 # Ask R how many unique elements are in glop$Dataset by nesting unique() inside length().
829 length (unique (glop$Dataset))
830
831 # >> And don't forget the other useful ones from Part 6.
832 # summary(), nrow(), ncol()

```

```

833 summary (glop)
834 nrow (glop)
835 ncol (glop)
836
837
838 #=====
839 # Part 10: Packages
840 #=====
841
842 # Packages are one of the features that makes R so powerful. Packages are basically new
843 # collections of functions that can be added into your R library of functions. Anybody can
844 # make a package, and when it is approved by the powers that be, it is distributed by the
845 # the Comprehensive R Archive Network (CRAN). As I write this, there are 6218 available
846 # packages.
847
848 # The most important packages are those for particular statistical analyses, and graphing.
849
850 # Functions are the vocabulary of the R language. Packages, therefore, are like dialects
851 # of R. It is easier to master a single dialect than many, and it is easiest to read and
852 # evaluate somebody else's code if it is in a dialect that you already know. I therefore
853 # highly recommend that you adopt new packages sparingly, sticking to the base-R language
854 # as much as possible.
855
856 # >> Installing a new package:
857 # Let's install ggplot2, the current standard for pretty graphics in R.
858 # First, set your 'CRAN mirror' - the place from which to download things.
859 # - Go to RStudio Options (or Preferences).
860 # - Click the Packages icon.
861 # - At the top, by 'CRAN mirror', click 'Change...'
862 # - A commonly used and reliable mirror is the USA (CA2) UC Berkeley mirror. Choose that.
863 # - Now, in the main dropdown menu, go to Tools --> Install packages...
864 # - 'Install from' = Repository (CRAN)
865 # - In the Packages text box, type in ggplot2 (an auto complete should show you package
866 # names as you type).
867 # - Make sure the 'Install dependencies' box is checked (many packages call on other
868 # packages).
869 # - Click 'Install'.
870
871 # >> Loading a package into R:
872 # You can have many packages installed on your computer, but their functions are not
873 # available until you load them into R.
874
875 #--Load ggplot2 using the library() function.
876 library (ggplot2)
877
878 #--Now look up help on the ggplot() function. Since ggplot2 is loaded, all the help
879 # documentation is available in the same format as the base-R functions.
880 ?ggplot
881
882 #--You can also use the require() function. This one can be better for use inside of a
883 # custom function. See the help documentation for the details.
884 require (ggplot2)
885
886
887 #=====
888 # Part 11: Basic troubleshooting
889 #=====
890
891 # THESE ARE DRILLS: Try only looking at the 'BLANK' script version and solving them.
892
893 # Troubleshooting is an art that you will develop with practice. Here are some basics.
894
895 #--Spelling: R will not do any 'fuzzy lookups' for you and suggest alternatives if you
896 # misspell something. So be careful, read carefully.

```

```

897 v <- vector.one
898 # Our vector.one had a period in it.
899 v <- vector.one
900
901 # >> Commas and parentheses:
902 # The most common cause of cryptic errors is misplacement of commas and parentheses.
903
904 # R will usually not know exactly what you should have done, it just tells you what
905 # doesn't make sense to it. Often, your problem is somewhere around where the red text in
906 # the error message ends.
907
908 #--What's wrong with this line?
909 m <- matrix (c(2,3,4)ncol=3)
910
911 # The corrected code:
912 m <- matrix (c(2,3,4), ncol=3)
913
914 #--How about this one? Attempting to paste the letter "a" to the end of each of entries 2
915 # through 6 and also 9 in glop.
916 glop[c(2:6,9),"Code"]=paste0(glop[c(2:6,9),"Code"],"a")
917
918 # Why is the ']' unexpected? What happens when you fix that and try again?
919
920 # The corrected code:
921 glop [c(2:6,9), "Code"] = paste0 (glop [c(2:6,9), "Code"], "a")
922
923 #--Note the usefulness of clear spacing! Imagine solving the above problem with this line
924 # of code. Much easier!
925 glop [c(2:6,9, "Code")] = paste0 (glop [c(2:6,9), "Code"] "a")
926
927 #=====
928 # POP QUIZ ANSWERS
929 #=====
930
931 # POP QUIZ 1: What happens if you DON'T put the quotes around "a"? Why did that happen?
932 # - Without quotes, R will look for an object named 'a' instead of the letter "a". If that
933 # object doesn't exist, R will return an error because it can't find 'a'.
934
935 # POP QUIZ 2: What unit is considered a single 'element' of df1? (Go back to the definition
936 # of data frames at the beginning of this Part.)
937 # - Data frames are 'lists' of 'vectors'. In other words, columns are vectors of equal
938 # length, and data frames are collections of columns. So a single 'element' of df1 is a
939 # column.
940
941
942

```